

CHALMERS UNIVERSITY OF TECHNOLOGY

FFR135 - Artificial Neural Networks

Examples sheet 4

Jacopo Credi
(910216-T396)

October 20, 2015

Task 1a




Figure 1: Energy (in log scale) vs iterations of backpropagation algorithm, for $N_G = 5$ Gaussian nodes. Training data in blue, validation data in red. For clarity, only three runs out of 20 are shown, each with only one point every 100 iterations. Parameters in the competitive learning phase: 10^5 iterations, learning rate $\eta_c = 0.02$, neighbourhood function width $\sigma = 0.1$. Parameters in the backpropagation phase: 10^4 iterations, learning rate $\eta_b = 0.1$, activation function $g(b) = \tanh(\beta b)$, with $\beta = 0.5$.

The provided data were classified using a neural network which combines unsupervised (competitive learning) and supervised learning (simple perceptron). For the competitive learning part, $N_G = 5$ Gaussian nodes were used. Given an input vector¹ \mathbf{x}^μ , the winning neuron (denoted by i_0) is the unit whose activation function is maximum: $g_{i_0}(\mathbf{x}^\mu) \geq g_i(\mathbf{x}^\mu)$ for all $i = 1, \dots, N_G$. After determining the winning unit, all weights are updated, according to the following learning rule:

$$\delta \mathbf{w}_i = \eta \Lambda(i, i_0) (\mathbf{x}^\mu - \mathbf{w}_i),$$

where the neighbouring function $\Lambda(i, i_0)$ was taken to be $\Lambda(i, i_0) = \exp[-||i - i_0||^2 / (2\sigma^2)]$. Then, the Gaussian nodes are used as input units of a simple perceptron, trained with backpropagation. In order to do this, the data were randomly divided in two parts: one for training (70%) and the other for validation (the remaining 30%).

Figure 1 shows the energy change as the backpropagation training is carried out. The energy quickly drops by a factor of ~ 8 in the first 2000 iterations, but then does not decrease further. The average final energy values, with parameters in figure caption and $N_G = 5$ Gaussian nodes, were

$$H^{(\text{training})} = (14 \pm 2) \cdot 10^1, \quad \text{and} \quad H^{(\text{validation})} = (6.0 \pm 1.1) \cdot 10^1$$

Values are averages over 20 runs, with standard deviation. The ratio of the values reflect the size of the two datasets, as energy is extensive. We will later see (1b) that such values correspond to a rather poor classification performance.

¹In this case, the data need not be normalised because we are not going to pass them (directly) through a sigmoid function, and there is no risk of zero derivative as in the perceptron case.

Task 1b




Figure 2: Classification error vs iterations of backpropagation algorithm, for $N_G = 5$ Gaussian nodes. Training data in blue, validation data in red. For clarity, only three runs out of 20 are shown, each with only one point every 100 iterations. Parameters in the competitive learning phase: 10^5 iterations, learning rate $\eta_c = 0.02$, neighbourhood function width $\sigma = 0.1$. Parameters in the backpropagation phase: 10^4 iterations, learning rate $\eta_b = 0.1$, activation function $g(b) = \tanh(\beta b)$, with $\beta = 0.5$.

The classification error can be computed as usual:

$$C_V = \frac{1}{2N_p} \sum_{\mu=1}^p | \zeta^{(\mu)} - \text{sgn}(O^{(\mu)}) | ,$$

where p is the number of data points (in this case, $p^{(\text{training})} = 1400$ and $p^{(\text{validation})} = 600$).

Using the same parameters as in **1a**, the classification error was computed every 100 backpropagation iterations, obtaining the plots in Fig. 2 (only 3 runs out of 20 are shown, for clarity).

The final classification errors for the training and validation data sets, with parameters in figure caption and $N_G = 5$ Gaussian nodes, were

$$C_V^{(\text{training})} = 0.069 \pm 0.010, \quad \text{and} \quad C_V^{(\text{validation})} = 0.070 \pm 0.014 .$$

Above values are averages over 20 runs, with standard deviation.

Although the training process can be considered complete (since both the energy and the error have reached their equilibrium value), the network “only” classifies 93% of the input points correctly. We will later see that the number of Gaussian nodes has a huge impact of the classification performance.

Task 1c




Figure 3: Left panel: data points, plotted in different colours and with different symbols according to the sign of the corresponding class, weight vectors as red crosses, and decision boundary as red line. Right panel: pseudo-colour map representation of the network output for a grid of equally spaced points covering the same portion of the input space. Same parameters as in **1a** and **1b**, 5 Gaussian nodes.

From the runs in **1a** and **1b**, the one with the best performance was chosen, i.e. the one yielding the lowest validation error (which turned out to also correspond to the lowest training error). Then, a large number (10^5) of random data points were generated and fed into the neural network. For each such point, if the output was in interval $[-0.01, 0.01]$, then the point was saved, otherwise it was discarded. The set of points obtained in this way was then sorted with respect to the x-coordinate, and then plotted, to produce the red line in the left panel of Fig. 3. In the same panel, the original data points are also plotted (in different colour, according to the corresponding class) and the Gaussian nodes weight vectors are shown as red crosses.

We can see that the weights correctly match the data distribution, but the decision boundary is not very precise and it incorrectly classifies a bunch of data points in both clusters. It can therefore be inferred that the number of Gaussian nodes is insufficient to fully capture the distribution of the data and classify them correctly.

The pseudo-colour map in the right panel of Fig. 3, instead, was obtained by generating a grid of equally spaced points in the input plane, passing these points through the network and colouring each point according to the network output (see colorbar).

Task 2a




Figure 4: Classification error vs iterations of backpropagation algorithm, for $N_G = 20$ Gaussian nodes. Training data in blue, validation data in red. For clarity, only three runs out of 20 are shown, each with only one point every 100 iterations. Parameters in the competitive learning phase: 10^5 iterations, learning rate $\eta_c = 0.02$, neighbourhood function width $\sigma = 0.1$. Parameters in the backpropagation phase: 10^4 iterations, learning rate $\eta_b = 0.1$, activation function $g(b) = \tanh(\beta b)$, with $\beta = 0.5$.

The network architecture was then modified, increasing the number of Gaussian nodes to 20 units. As shown in Fig. 4, this led to significantly smaller classification errors, or in other terms, to a better classification performance.

The final classification errors for the training and validation data sets, with parameters in figure caption and $N_G = 20$ Gaussian nodes, were in fact

$$C_V^{(\text{training})} = 0.001 \pm 0.001, \quad \text{and} \quad C_V^{(\text{validation})} = 0.001 \pm 0.002 .$$

Above values are averages over 20 runs, with standard deviation. Indeed, values are statistically compatible with zero. The network is correctly classifying input data 99.9% of the time.

Task 2b




Figure 5: Left panel: data points, plotted in different colours and with different symbols according to the sign of the corresponding class, weight vectors as red crosses, and decision boundary as red line. Right panel: pseudo-colour map representation of the network output for a grid of equally spaced points covering the same portion of the input space. Same parameters as in **2a**, 20 Gaussian nodes.

From the runs in **2a**, the one with the best performance was chosen, i.e. the one yielding the lowest validation error (which again turned out to also correspond to the lowest training error). Then, a large number (10^5) of random data points were generated and fed into the neural network. For each such point, if the output was in interval $[-0.01, 0.01]$, then the point was saved, otherwise it was discarded. The set of points obtained in this way was then sorted with respect to the x-coordinate, and then plotted, to produce the red line in the left panel of Fig. 5. In the same panel, the original data points are also plotted (in different colour, according to the corresponding class) and the Gaussian nodes weight vectors are shown as red crosses.

Again, the weights correctly match the data distribution, spreading roughly uniformly in the space covered by the two data clusters. Compared to the one obtained in **1c**, the decision boundary separates the data clusters much more accurately. It can therefore be inferred that 20 Gaussian nodes are sufficient to fully capture the distribution of the data and classify them correctly.

Task 3




Figure 6: Classification error vs number of Gaussian nodes used for the competitive learning part. Training data in blue, validation data in red. Parameters in the competitive learning phase: 10^5 iterations, learning rate $\eta_c = 0.02$, neighbourhood function width $\sigma = 0.1$. Parameters in the backpropagation phase: $5 \cdot 10^3$ iterations, learning rate $\eta_b = 0.1$, activation function $g(b) = \tanh(\beta b)$, with $\beta = 0.5$.

The effect of the number of Gaussian nodes on the network performance was tested by using a variable number N_G of nodes $N_G = 1, 2, \dots, 20$. For each value of the number of Gaussian nodes, 20 independent training runs were carried out.

The resulting final average classification error as a function of the number of nodes is plotted in Figure 6. Errorbars are standard deviations over the 20 runs. Errors are displayed for both the (randomly constructed) training (in blue) and validation set (in red), although points exhibit a strong overlap and are hard to distinguish.

There are two main remarks about this plot. First, with a single Gaussian node, the error is about 0.5, i.e the network is correctly classifying input points 50% of the time, just like a random classifier would do. This is expected, since data are organised in two clusters, and this feature cannot be captured by using a single Gaussian node in the competitive unsupervised learning phase. Two Gaussian nodes, however, are still too few for a good classification, as clusters are not “convex” and several data points belonging to one of them are actually closer² to the “centre” of the *other* cluster (the one that they do *not* belong to).

A number of Gaussian nodes of about $9 \sim 10$ appears to be sufficient to properly reconstruct the distribution of the original data and therefore train the network as a classifier, provided that the training is carried out with suitable parameters and for a sufficient number of iterations.

²“Closer” here refers to both the Euclidean and the radial basis function sense.

MATLAB code used

ClassificationProblem.m

```
1 %% CLASSIFICATIONPROBLEM Main script for ANN - Examples sheet 4
2 clear; clc; close all;
3
4 %% Parameters
5 nRuns = 20;
6 nCurvesToPlot = 3;
7
8 nIterationsKohonen = 1e5;
9 nKohonenNodes = 20;
10 etaKohonen = 0.02;
11 sigmaNeighbourhoodFunction = 0.2;
12
13 nIterationsBackpropagation = 5e3;
14 nOutputNodes = 1;
15 etaBackpropagation = 0.1;
16 beta = 0.5;
17 nIterationsBetweenPlots = 100;
18
19 %% Load data
20 tmpData = importdata(' ../Data/data_classify.txt ');
21 patternClasses = tmpData(:,1);
22 patterns = tmpData(:,2:end);
23 nPatterns = size(patterns, 1);
24
25 %% Loop over number of Gaussian nodes
26 % h = waitbar(0,'Please wait...');
27 % for nKohonenNodes=1:20
28
29 %% Loop over runs
30 for iRun = 1:nRuns
31
32     fprintf('\nRUN %i:\n',iRun);
33
34     % Unsupervised learning part (Kohonen)
35     tic;
36     kohonenWeights = CompetitiveLearning(patterns, nKohonenNodes, ...
37         sigmaNeighbourhoodFunction, nIterationsKohonen, etaKohonen);
38     fprintf(' - Competitive learning completed in %4.3f seconds.\n',toc);
39
40     % Supervised learning part (backpropagation)
41     tic;
42     resultsStructure(iRun) = Backpropagation(patterns,patternClasses,
43         kohonenWeights, nIterationsBackpropagation, beta,
44         etaBackpropagation, nIterationsBetweenPlots);
```



```
43     fprintf(' - Perceptron backpropagation completed in %4.3f seconds.\n',  
44             ,toc);  
45 end  
46 clc;  
47 fprintf('%i runs with %i Kohonen nodes completed.\n',nRuns,nKohonenNodes  
48         );  
49 % Plot energy and error (3 curves)  
50 whichOnes = randperm(nRuns,nCurvesToPlot);  
51 [energyTrainingPlot, energyValidationPlot] = PlotEnergy(whichOnes, ...  
52     nIterationsBetweenPlots, nIterationsBackpropagation, resultsStructure)  
53     ;  
54 [errorTrainingPlot, errorValidationPlot] = PlotError(whichOnes,...  
55     nIterationsBetweenPlots, nIterationsBackpropagation, resultsStructure)  
56     ;  
57 % Average energy and error (with uncertainties) for training and  
58 % validation  
59 tmpEnTr = zeros(nRuns,1);  
60 tmpEnVal = zeros(nRuns,1);  
61 tmpErrTr = zeros(nRuns,1);  
62 tmpErrVal = zeros(nRuns,1);  
63 for iRun = 1:nRuns  
64     tmpEnTr(iRun) = resultsStructure(iRun).energyTraining(end);  
65     tmpEnVal(iRun) = resultsStructure(iRun).energyValidation(end);  
66     tmpErrTr(iRun) = resultsStructure(iRun).errorTraining(end);  
67     tmpErrVal(iRun) = resultsStructure(iRun).errorValidation(end);  
68 end  
69 avgEnergyTraining = mean(tmpEnTr);  
70 errEnergyTraining = std(tmpEnTr);  
71 avgEnergyValidation = mean(tmpEnVal);  
72 errEnergyValidation = std(tmpEnVal);  
73 avgErrorTraining = mean(tmpErrTr);  
74 errErrorTraining = std(tmpErrTr);  
75 avgErrorValidation = mean(tmpErrVal);  
76 errErrorValidation = std(tmpErrVal);  
77 [~,iBestClassifier] = min(tmpErrVal);  
78 % errVsNodesTrainingAverage(nKohonenNodes) = avgErrorTraining;  
79 % errVsNodesTrainingSd(nKohonenNodes) = errErrorTraining;  
80 % errVsNodesValidationAverage(nKohonenNodes) = avgErrorValidation;  
81 % errVsNodesValidationSd(nKohonenNodes) = errErrorValidation;  
82 fprintf('\n Final energy (training set): %4.3f, uncertainty: %4.3f',  
83     avgEnergyTraining,errEnergyTraining);  
84 fprintf('\n Final energy (validation set): %4.3f, uncertainty: %4.3f',  
85     avgEnergyValidation,errEnergyValidation);
```

```

85 fprintf('\n\n Final classification error (training set): %4.3f,
    uncertainty: %4.3f',avgErrorTraining,errErrorTraining);
86 fprintf('\n Final classification error (validation set): %4.3f,
    uncertainty: %4.3f\n\n',avgErrorValidation,errErrorValidation);
87
88 disp('Run script "DecisionBoundary.m" for a graphical representation of
    the decision boundary for the best network found.');
```

CompetitiveLearning.m

```

1 function kohonenWeights = CompetitiveLearning(patterns, nKohonenNodes,
    ...
2     sigmaNeighbourhoodFunction, nIterationsKohonen, etaKohonen)
3 %% CompetitiveLearning
4
5 nPatterns = size(patterns, 1);
6
7 % Initialise weights
8 kohonenWeights = -1 + 2*rand(nKohonenNodes,2);
9
10 % trick for speed
11 tmp = 1:nKohonenNodes;
12 A = repmat(tmp,nKohonenNodes,1)-repmat(tmp',1,nKohonenNodes);
13 A = triu(A) + triu(A)';
14 neighbourhoodMatrix = exp(-(A.^2)/(2*sigmaNeighbourhoodFunction^2));
15
16 % =====
17 %% Plot weights update on-the-go (cool, but not recommended for nRuns >
    1)
18 % figure();
19 % hold on
20 % dataPlot = plot(patterns(:,1),patterns(:,2),'.','MarkerSize',12);
21 % weightsPlot = plot(kohonenWeights(:,1),kohonenWeights(:,2),'x');
22 % set(weightsPlot,'LineWidth',2,'MarkerSize',20);
23 % hold off
24 % =====
25
26 %% Training loop
27 for iKohonen = 1:nIterationsKohonen
28
29     iPattern = randi(nPatterns);
30     inputPattern = patterns(iPattern,:);
```

```

31
32     iWinning = GetWinningNeuron(inputPattern, kohonenWeights);
33     kohonenWeights = UpdateWeights(kohonenWeights, iWinning, inputPattern
34         ,...
35         etaKohonen, neighbourhoodMatrix);
36
37 % =====
38 %     if mod(iKohonen,1000) == 0
39 %         set(weightsPlot,'XData',kohonenWeights(:,1));
40 %         set(weightsPlot,'YData',kohonenWeights(:,2));
41 %         drawnow;
42 %     end
43 % =====
44 end
45
end

```

GetWinningNeuron.m

```

1 function iWinning = GetWinningNeuron(inputPattern, kohonenWeights)
2
3 nNodes = size(kohonenWeights,1);
4 activationFunction = zeros(nNodes,1);
5
6 for j = 1:nNodes
7     distance = sqrt((inputPattern(1)-kohonenWeights(j,1))^2 + ...
8         (inputPattern(2)-kohonenWeights(j,2))^2);
9     activationFunction(j) = exp(-distance/2);
10 end
11
12 [~,iWinning] = max(activationFunction);
13
14 end

```

UpdateWeights.m

```

1 function updatedWeights = UpdateWeights(weights, iWinning, inputPattern
2     ,...
3     etaKohonen, neighbourhoodMatrix)
4
5 nNodes = size(weights,1);
6 deltaWeights = zeros(nNodes,2);
7
8 for iNode = 1:nNodes
9     tmpDiff = inputPattern-weights(iNode,:);

```

```

9      deltaWeights(iNode,:) = etaKohonen*neighbourhoodMatrix(iNode,iWinning)
      .*tmpDiff;
10 end
11
12 updatedWeights = weights + deltaWeights;
13
14 end

```

Backpropagation.m

```

1 function resultsStructure = Backpropagation(patterns, patternClasses,
      kohonenWeights, ...
2      nIterationsBackpropagation, beta, etaBackpropagation,
      nIterationsBetweenPlots)
3 %% Backpropagation
4
5 nPatterns = size(patterns, 1);
6 nKohonenNodes = size(kohonenWeights,1);
7 nPlotPoints = 1+fix(nIterationsBackpropagation/nIterationsBetweenPlots);
8
9 % Randomly split data in training and validation sets
10 nTrainingData = fix(0.7*nPatterns);
11 trainingData = randperm(nPatterns, nTrainingData);
12 validationData = setdiff(1:nPatterns, trainingData);
13
14 % Initialisations
15 perceptronWeights = -1 + 2*rand(1,nKohonenNodes+1);
16 energyTraining = zeros(1, nPlotPoints);
17 errorTraining = zeros(1, nPlotPoints);
18 energyValidation = zeros(1, nPlotPoints);
19 errorValidation = zeros(1, nPlotPoints);
20 iPlot = 1;
21
22 % Compute initial energy and error
23 [energyTraining(iPlot), errorTraining(iPlot)] = ComputeEnergyAndError
      (...
24 patterns, patternClasses, kohonenWeights, perceptronWeights, beta);
25 [energyValidation(iPlot), errorValidation(iPlot)] =
      ComputeEnergyAndError(...
26 patterns, patternClasses, kohonenWeights, perceptronWeights, beta);
27
28 %=====
29 %% Draw colormap on-the-go (cool, but not recommended for nRuns > 1)
30 %
31 % rangeX = [-15, 25];
32 % rangeY = [-10, 15];
33 % spacing = 0.33;

```

```

34 % [randomDataX,randomDataY] = meshgrid(rangeX(1):spacing:rangeX(2),
    rangeY(1):spacing:rangeY(2));
35 % nY = size(randomDataX,1);
36 % nX = size(randomDataX,2);
37 % randomDataZ = zeros(size(randomDataX));
38 % % Pass through classifier
39 % for iY = 1:nY
40 %     for iX = 1:nX
41 %         inputPattern = [randomDataX(iY,iX), randomDataY(iY,iX)];
42 %         [~, randomDataZ(iY,iX)] = RunClassifier(inputPattern, ...
43 %             kohonenWeights, nKohonenNodes, perceptronWeights, beta);
44 %     end
45 % end
46 % % Pseudocolor map
47 % pcolor(randomDataX,randomDataY,randomDataZ)
48 % shading flat
49 % axis square;
50 %=====
51
52 for iBackpropagation = 1:nIterationsBackpropagation
53
54     % pick pattern from training set
55     iPattern = trainingData(randi(nTrainingData));
56     inputPattern = patterns(iPattern,:);
57     desiredOutput = patternClasses(iPattern);
58
59     % run classifier
60     [gaussianActivation, output] = RunClassifier(inputPattern, ...
61         kohonenWeights, nKohonenNodes, perceptronWeights, beta);
62
63     % backpropagation
64     gPrime = beta*(1-output.^2); % because we're using g = tanh
65     outputError = desiredOutput - output;
66     deltaWeights = etaBackpropagation*outputError*gPrime*
        gaussianActivation';
67     perceptronWeights = perceptronWeights + deltaWeights;
68
69     % compute energy and classification error
70     if mod(iBackpropagation,nIterationsBetweenPlots) == 0
71         iPlot = iPlot + 1;
72         [energyTraining(iPlot), errorTraining(iPlot)] =
            ComputeEnergyAndError(...
73             patterns(trainingData,:), patternClasses(trainingData), ...
74             kohonenWeights, perceptronWeights, beta);
75         [energyValidation(iPlot), errorValidation(iPlot)] =
            ComputeEnergyAndError(...
76             patterns(validationData,:), patternClasses(validationData), ...
77             kohonenWeights, perceptronWeights, beta);
78

```

```

79 %=====
80 % % Pass random data through classifier
81 % for iY = 1:nY
82 %     for iX = 1:nX
83 %         inputPattern = [randomDataX(iY,iX), randomDataY(iY,iX)];
84 %         [~, randomDataZ(iY,iX)] = RunClassifier(inputPattern, ...
85 %             kohonenWeights, nKohonenNodes, perceptronWeights, beta);
86 %     end
87 % end
88 % % Pseudocolor map
89 % pcolor(randomDataX,randomDataY,randomDataZ);
90 % shading flat;
91 % axis square;
92 % pause(0.001);
93 %=====
94 end
95 end
96
97 % Save results in a structure
98 resultsStructure.kohonenWeights = kohonenWeights;
99 resultsStructure.perceptronWeights = perceptronWeights;
100 resultsStructure.energyTraining = energyTraining;
101 resultsStructure.energyValidation = energyValidation;
102 resultsStructure.errorTraining = errorTraining;
103 resultsStructure.errorValidation = errorValidation;
104
105 end

```

RunClassifier.m

```

1 function [gaussianActivation, output] = RunClassifier(inputPattern, ...
2     kohonenWeights, nKohonenNodes, perceptronWeights, beta)
3 %% RunClassifier
4
5 % run Kohonen network
6 gaussianActivation = zeros(nKohonenNodes+1,1);
7 for j = 1:nKohonenNodes
8     distance = sqrt((inputPattern(1)-kohonenWeights(j,1))^2 + ...
9         (inputPattern(2)-kohonenWeights(j,2))^2);
10    gaussianActivation(j) = exp(-distance/2);
11 end
12 gaussianActivation = gaussianActivation./sum(gaussianActivation); %
    normalise
13 gaussianActivation(end) = -1; % dummy neuron (for threshold)
14
15 % run perceptron
16 output = tanh(beta*(perceptronWeights*gaussianActivation));

```

```

17
18 end

```

ComputeEnergyAndError.m

```

1 function [energy, classificationError] = ComputeEnergyAndError(patterns
   ,...
2   patternClasses, kohonenWeights, perceptronWeights, beta)
3 %% ComputeEnergyAndError
4
5 nPatterns = size(patterns, 1);
6 nKohonenNodes = size(kohonenWeights,1);
7
8 tmpOutput = zeros(nPatterns,1);
9
10 for iPattern = 1:nPatterns
11   inputPattern = patterns(iPattern,:);
12
13   [~, tmpOutput(iPattern)] = RunClassifier(inputPattern, kohonenWeights,
   ...
14   nKohonenNodes, perceptronWeights, beta);
15 end
16
17 % compute energy
18 tmpEnergy = (patternClasses - tmpOutput).^2;
19 energy = 0.5*sum(tmpEnergy);
20
21 % compute classification error
22 tmpError = abs(patternClasses - sign(tmpOutput));
23 classificationError = 0.5/nPatterns*sum(tmpError);
24
25 end

```

PlotEnergy.m

```

1 function [energyTrainingPlot, energyValidationPlot] = PlotEnergy(
   whichOnes,...
2   nIterationsBetweenPlots, nIterationsBackpropagation, resultsStructure)
3 %% PlotEnergy
4
5 xValues = 1:nIterationsBetweenPlots:(nIterationsBackpropagation+1);
6 nCurvesToPlot = numel(whichOnes);
7
8 parulaColours = get(groot, 'DefaultAxesColorOrder');
9 defaultBlue = parulaColours(1,:);
10 defaultRed = parulaColours(2,:);

```

```

11
12 figure('Units','normalized','OuterPosition',[0.15 0.15 0.7 0.7]);
13 set(gcf, 'Color','w');
14 hold on
15
16 for i = 1:nCurvesToPlot
17
18     iPlot = whichOnes(i);
19
20     % training
21     energyTrainingPlot(iPlot) = semilogy(xValues,resultsStructure(iPlot).
        energyTraining);
22     set(energyTrainingPlot(iPlot),'Color',defaultBlue,'LineWidth',1.5);
23
24     % validation
25     energyValidationPlot(iPlot) = semilogy(xValues,resultsStructure(iPlot)
        .energyValidation);
26     set(energyValidationPlot(iPlot),'Color',defaultRed,'LineWidth',1.5);
27
28 end
29
30 set(gca,'yscale','log','FontSize',16);
31 xlabel('Backpropagation iterations');
32 ylabel('Energy H');
33 xlim([0,nIterationsBackpropagation]);
34 set(0, 'DefaultAxesBox', 'on');
35 legend([energyTrainingPlot(iPlot), energyValidationPlot(iPlot)], '
        Training data','Validation data');
36 pbaspect([1.618 1 1])
37 hold off
38
39 end

```

PlotError.m

```

1 function [errorTrainingPlot, errorValidationPlot] = PlotError(whichOnes
    ,...
2     nIterationsBetweenPlots, nIterationsBackpropagation, resultsStructure)
3 %% PlotError
4
5 xValues = 1:nIterationsBetweenPlots:(nIterationsBackpropagation+1);
6 nCurvesToPlot = numel(whichOnes);
7
8 parulaColours = get(groot,'DefaultAxesColorOrder');
9 defaultBlue = parulaColours(1,:);
10 defaultRed = parulaColours(2,:);
11

```



```

12 figure('Units','normalized','OuterPosition',[0.15 0.15 0.7 0.7]);
13 set(gcf,'Color','w');
14 hold on
15
16 for i = 1:nCurvesToPlot
17
18     iPlot = whichOnes(i);
19
20     % training
21     errorTrainingPlot(iPlot) = plot(xValues,resultsStructure(iPlot).
22         errorTraining);
23
24     % validation
25     errorValidationPlot(iPlot) = plot(xValues,resultsStructure(iPlot).
26         errorValidation);
27
28     set(errorTrainingPlot(iPlot),'Color',defaultBlue,'LineWidth',1.5);
29
30     set(errorValidationPlot(iPlot),'Color',defaultRed,'LineWidth',1.5);
31
32 end
33
34 set(gca,'FontSize',16);
35 xlabel('Backpropagation iterations');
36 ylabel('Classification error');
37 xlim([0,nIterationsBackpropagation]);
38 set(0,'DefaultAxesBox','on');
39 legend([errorTrainingPlot(iPlot), errorValidationPlot(iPlot)],'Training
40     data','Validation data');
41 pbaspect([1.618 1 1])
42 hold off
43
44 end

```

DecisionBoundary.m

```

1 %% DECISIONBOUNDARY Script to draw decision boundary and colormap
2
3 nRandomData = 1e5;
4 rangeX = [-15, 25];
5 rangeY = [-10, 15];
6 spacing = 0.5;
7 beta = 0.5;
8 threshold = 0.01;
9
10 bestKohonenWeights = resultsStructure(iBestClassifier).kohonenWeights;
11 bestPerceptronWeights = resultsStructure(iBestClassifier).
12     perceptronWeights;

```

```

13 %% Initialise plot
14 figure('Units','normalized','OuterPosition',[0.15 0.15 0.7 0.7]);
15 set(gcf, 'Color','w');
16 hold on
17 parulaColours = get(groot,'DefaultAxesColorOrder');
18 colour1 = parulaColours(1,:);
19 colour2 = parulaColours(3,:);
20
21 %% Generate random data (for decision boundary)
22 randomDataX = rangeX(1) + diff(rangeX)*rand(nRandomData,1);
23 randomDataY = rangeY(1) + diff(rangeY)*rand(nRandomData,1);
24 randomDataZ = zeros(nRandomData,1);
25 % Pass through classifier
26 for iRandomData = 1:nRandomData
27     inputPattern = [randomDataX(iRandomData), randomDataY(iRandomData)];
28     [~, randomDataZ(iRandomData)] = RunClassifier(inputPattern,
29         bestKohonenWeights, ...
30         nKohonenNodes, bestPerceptronWeights, beta);
31 end
32 % Generate decision boundary
33 decisionBoundaryX = randomDataX(abs(randomDataZ) < threshold);
34 decisionBoundaryY = randomDataY(abs(randomDataZ) < threshold);
35 [decisionBoundaryX, order] = sort(decisionBoundaryX);
36 decisionBoundaryY = decisionBoundaryY(order);
37
38 % Subplot 1
39 subplot(1,2,1);
40 plotCluster1 = plot(patterns(1:1000,1),patterns(1:1000,2));
41 set(plotCluster1,'LineStyle','none','Marker','s','MarkerSize',6,'
42     LineWidth',1,'Color',colour2);
43 hold on
44
45 plotCluster2 = plot(patterns(1001:end,1),patterns(1001:end,2),'o','
46     MarkerSize',10);
47 set(plotCluster2,'LineStyle','none','Marker','o','MarkerSize',6,'
48     LineWidth',1,'Color',colour1);
49
50 plotWeights = plot(bestKohonenWeights(:,1),bestKohonenWeights(:,2),'x');
51 set(plotWeights,'LineWidth',2,'MarkerSize',16,'Color','r');
52
53 plotDecisionBoundary = plot(decisionBoundaryX, decisionBoundaryY,'r','
54     LineWidth',2);
55 set(gca,'FontSize',16);
56 axis square;
57 xlim([-15, 25]);
58 hsp1 = get(gca, 'Position');
59 hold off
60
61 %% Generate meshgrid (for pseudo-colour map)

```

```
57 [meshGridX,meshGridY] = meshgrid(rangeX(1):spacing:rangeX(2), rangeY(1):  
    spacing:rangeY(2));  
58 nY = size(meshGridX,1);  
59 nX = size(meshGridX,2);  
60 meshGridZ = zeros(size(meshGridX));  
61 % Pass through classifier  
62 for iY = 1:nY  
63     for iX = 1:nX  
64         inputPattern = [meshGridX(iY,iX), meshGridY(iY,iX)];  
65         [~, meshGridZ(iY,iX)] = RunClassifier(inputPattern, ...  
66             bestKohonenWeights, nKohonenNodes, bestPerceptronWeights, beta);  
67     end  
68 end  
69  
70 % Plot pseudo-colour map  
71 subplot(1,2,2);  
72 pcolor(meshGridX,meshGridY,meshGridZ)  
73 shading flat  
74 axis square;  
75 hsp2 = get(gca, 'Position');  
76 colorbar;  
77 set(gca, 'FontSize',16);  
78 set(gca, 'Position', [hsp2(1:2) hsp1(3:4)]);
```